

Table of Contents

1	Introduction	1
2	Background & Related Work	3
2.1	Introduction to NVIDIA Ampere Architecture	3
2.2	Transformers	4
2.2.1	T5	6
2.2.2	NLLB	7
2.2.3	NLLB-MoE & Switch Transformers	7
2.3	Performance Evaluation on MT models	8
2.3.1	Quality Metrics	8
2.3.2	Model Inference Time	11
2.4	Disk offloading	12
2.4.1	Related work: Deepspeed	12
2.4.2	Archer	13
3	Benchmark Design and Implementation	15
3.1	BenchmarkingSuite	15
3.2	Configuration Setup	17
3.2.1	ExperimentConfiguration	17
3.3	Dataset Class	17
3.3.1	Archer	18
3.4	Analytics	18
3.4.1	Quality Metrics	19
3.4.2	Efficiency Metrics	20
3.4.3	Profiling	20
3.5	Implementation	20

4	Experiments	22
4.1	Datasets & Models	22
4.2	Hyperparameters	25
4.3	Methodology	27
4.4	Results & Analysis	29
4.4.1	Dense Models	29
4.4.2	Sparse Models	34
5	Conclusions	39
	Bibliography	41
6	Appendix	47
6.1	Single sentence dataset from FLORES-200 for Examining Efficiency of Models	47
6.2	Single sentence dataset from WMT14 for Examining Efficiency of Models	47
6.3	Training hyperparameters of the SwitchTransformer Model Fine-Tuned	48